

PRD: Magnifica Humanitas Site Upgrade

Turning the Encyclical from a Passive Document into an Operational Doctrine Engine

Conclusion / Summary

The current site succeeds as a solemn encyclical-style artifact, but it needs a stronger UX layer that helps readers **understand, navigate, extract, and apply** the doctrine.

This PRD instructs the coding agent to add:

- Three reading modes
- Principle cards
- Governance gates
- Doctrine map
- Reader paths
- "Five Red Lines" section
- "Person Test" section
- Quote actions
- Stronger chapter headers
- A governance / ethics export area
- Better mobile and accessibility behavior

The goal is not to make the site feel like a SaaS dashboard. The goal is to make it feel like a **cathedral-shaped AI governance instrument**.

Traditional Approach:

1. Product Objective

Upgrade the **Magnifica Humanitas** site from a long-form encyclical page into a more usable, navigable, and actionable governance experience.

The site should still preserve its ceremonial, theological, and ecclesial tone, but it should also help users:

- Understand the encyclical quickly
- Extract the key doctrine
- Apply the doctrine to AI systems
- Copy/use principles as governance rules

- Navigate the long document without fatigue
 - See the moral argument as a system, not just a text
-

2. Current Problem

The current page has strong content and atmosphere, but the UX is too passive.

Observed issues:

- The site is mostly linear and long-form
 - Strong ideas are buried inside dense sections
 - No clear “use this doctrine” pathway
 - No applied governance layer
 - No principle-to-action translation
 - Pull quotes are not operationalized
 - Readers are not guided by role or purpose
 - The page lacks an executive summary layer
 - There is no exportable governance artifact
 - Mobile navigation may become heavy because of document length
-

3. Target Users

Primary Users

AI governance / ethics reviewers

They need principles, gates, risk checks, and decision criteria.

Builders / founders

They need to know what kinds of AI systems the doctrine allows, warns against, or blocks.

Policy and institutional readers

They need fast orientation, legitimacy, and application areas.

Theological / Catholic Social Doctrine readers

They need continuity with social doctrine and clarity of argument.

Creative technologists / public intellectuals

They need a strong conceptual artifact that can be quoted, reused, and translated into tools.

4. Core UX Strategy

Create three layers:

Layer 1 — The Encyclical

The full ceremonial text remains intact.

Purpose:

- Authority
- Beauty
- Depth
- Theological continuity
- Long-form reading

Layer 2 — The Doctrine Map

Extract and organize the moral framework.

Purpose:

- Scannability
- Memorability
- Conceptual clarity
- Principle-level navigation

Layer 3 — The Governance Engine

Turn principles into practical evaluation tools.

Purpose:

- AI system review
- Governance gates
- Red lines
- Actionable prompts
- Exportable agent skill

5. Feature Requirements

Feature 1: Add Executive Summary Block

Location

Immediately after the cover / opening hero and before the full table of contents.

Purpose

Give users a quick way to understand the central argument before reading the full encyclical.

Required Copy

Title:

A Doctrine for Judging Machines by What They Do to the Human Person

Body:

Magnifica Humanitas argues that artificial intelligence must be judged by whether it protects or diminishes the human person. Its central claim is simple: human dignity cannot be earned, measured, optimized, automated, or delegated away.

Supporting line:

The machine may calculate. It may not absolve.

UI Requirements

- Use same ceremonial visual language as the current site
- Must not feel like a startup hero section
- Should feel like an illuminated doctrinal preface
- Include subtle divider between summary and table of contents
- Responsive on mobile

Acceptance Criteria

- User can understand the core thesis within 10 seconds
 - Summary is visible without needing to scroll deeply
 - Visual style remains consistent with encyclical atmosphere
-

Feature 2: Add Reading Mode Selector

Location

Near top of page, after Executive Summary.

Modes

1. Read the Encyclical

Shows the full long-form document.

2. Extract the Doctrine

Jumps to doctrine map, principles, quotes, and red lines.

3. Apply the Doctrine

Jumps to governance gates, person test, and export area.

Behavior

- These can be anchor buttons, tabs, segmented controls, or cards
- Do not hide the full document unless existing architecture supports it cleanly
- Preferred MVP: anchor-based navigation
- Add smooth scroll behavior
- Preserve current dropdown navigation

Acceptance Criteria

- User can choose a reading path immediately
 - Buttons work on desktop and mobile
 - No content becomes inaccessible
-

Feature 3: Add “Five Principles for the Digital Age” Cards

Location

After Chapter II summary or before the main Chapter II body.

Principles

Create five cards:

1. Common Good

Doctrine:

Not the mere sum of interests, but a “plus” born of interdependence.

AI translation:

AI systems must be judged by their effects on shared life, public trust, access, and institutional power.

Governance question:

Who benefits, who pays the cost, and who is excluded?

Red flag:

A system concentrates power while distributing harm.

2. Universal Destination of Goods

Doctrine:

The goods of the earth belong to all — including patents, algorithms, platforms, infrastructure, and data.

AI translation:

Data and digital infrastructure must not become private empires when they shape public life.

Governance question:

Is this system enclosing a shared good?

Red flag:

Public data is harvested into private control without return, remedy, or accountability.

3. Subsidiarity

Doctrine:

Higher powers must not supplant lower communities.

AI translation:

AI must preserve local agency, human appeal, and community-level override.

Governance question:

Who can contest or override the system?

Red flag:

A centralized model imposes decisions with no meaningful appeal.

4. Solidarity

Doctrine:

Networks create nearness, but solidarity is the conscious choice to bear one another.

AI translation:

AI should not exploit invisible labor, vulnerable users, or distant communities.

Governance question:

Whose labor, data, attention, or environment makes this system possible?

Red flag:

The system hides the people who pay its real cost.

5. Social Justice

Doctrine:

Place the person, not wealth or data accumulation, at the center.

AI translation:

AI must not create new exclusions through surveillance, opacity, scoring, or automation.

Governance question:

Does this system make the vulnerable more vulnerable?

Red flag:

The system quietly punishes people who cannot be easily measured.

UI Requirements

- Use card layout
- Each card should have doctrine, AI translation, governance question, and red flag
- Cards should stack cleanly on mobile
- Avoid colorful SaaS styling
- Use ecclesial / archival / constitutional visual treatment

Acceptance Criteria

- All five cards are visible and readable
 - Cards can be used as standalone governance references
 - Mobile layout does not become cramped
-

Feature 4: Add Doctrine Map

Location

After the principle cards.

Purpose

Give readers a one-glance map of the entire encyclical.

Required Structure

Title:

Doctrine Map

Subtitle:

From theological foundation to AI governance action.

Sections:

Foundation

- Human dignity
- Common good
- Subsidiarity
- Solidarity
- Social justice

Threats

- Technocratic paradigm
- AI mystification
- Data enclosure
- Digital colonialism
- Automated war
- Synthetic falsehood

Protected Goods

- Truth
- Work
- Freedom
- Human weakness
- Peace
- Communion

Required Actions

- Transparent algorithms
- Worker protection
- Human oversight
- Data stewardship
- Digital sobriety
- Multilateral governance

UI Requirements

- Four-column desktop layout
- Two-column tablet layout

- One-column mobile layout
- Use clear visual hierarchy
- Consider subtle lines / borders to make it feel like a doctrinal chart

Acceptance Criteria

- User can understand the structure of the encyclical without reading the whole document
 - Doctrine map visually separates foundations, threats, protected goods, and actions
-

Feature 5: Add Governance Gates

Location

Create a new section before the conclusion or before Chapter V.

Title

Governance Gates

Subtitle

No AI system should pass without answering these questions.

Required Gates

Gate 1: Human Dignity Gate

Question:

Does the system reduce a person to a score, prediction, risk category, productivity measure, or behavioral value?

Pass condition:

The system preserves human dignity regardless of efficiency, productivity, disability, age, poverty, vulnerability, or predicted value.

Fail condition:

The system treats people primarily as costs, risks, resources, attention targets, or optimization variables.

Gate 2: Common Good Gate

Question:

Does the system serve shared life, or does it enclose power?

Pass condition:

Benefits and burdens are identified, distributed, and governed with public-interest safeguards.

Fail condition:

The system extracts value from many while concentrating control in a few.

Gate 3: Subsidiarity Gate

Question:

Can affected people contest, appeal, override, or meaningfully influence the system?

Pass condition:

The system includes appeal paths, human review, stakeholder input, and local/contextual override.

Fail condition:

Decisions are imposed by centralized technical, corporate, or state actors without redress.

Gate 4: Truth Gate

Question:

Does the system blur truth and falsehood, or does it disclose uncertainty and provenance?

Pass condition:

Synthetic content, uncertainty, sources, recommendation logic, and limits are clearly disclosed.

Fail condition:

The system presents opaque, manipulative, unverifiable, or synthetic output as trustworthy truth.

Gate 5: Accountability Gate

Question:

Can a harmed person identify who is responsible?

Pass condition:

There is a named human owner, audit trail, escalation path, appeal process, and final human authority.

Fail condition:

Responsibility is hidden behind the model, vendor, automation layer, or “the machine.”

UI Requirements

- Gates should feel serious and judicial
- Use expandable cards or stacked panels
- Each gate needs: question, pass, fail
- Optional: add small labels: “Question,” “Pass,” “Fail”
- Keep readable on mobile

Acceptance Criteria

- User can use the gates to review a real AI system
 - Each gate is visually distinct
 - Gate labels are accessible and readable
-

Feature 6: Add “Five Red Lines” Section

Location

After Governance Gates.

Purpose

Make the doctrine memorable and enforceable.

Required Copy

Title:

Five Red Lines

Subtitle:

What the doctrine refuses.

Red lines:

- No human being may be reduced to a score of worth.
- No high-stakes decision may hide responsibility behind a model.
- No private platform may quietly govern public truth without scrutiny.
- No automation may sacrifice workers without protection and retraining.
- No machine may simulate moral agency where conscience is required.

UI Requirements

- Should feel like a solemn prohibition list
- Use strong typography
- Consider numbered Roman numerals I-V
- Avoid aggressive red unless already supported by design system

Acceptance Criteria

- Red lines are immediately scannable
 - Section has visual weight
 - It feels doctrinal, not like a marketing list
-

Feature 7: Add “The Person Test”

Location

After Five Red Lines.

Purpose

Create a simple practical tool that any reader can apply.

Required Copy

Title:

The Person Test

Subtitle:

Before every system, ask who becomes more human — and who becomes easier to discard.

Questions:

- Who becomes more visible to power?
- Who becomes less able to appeal?
- Who carries the cost of efficiency?
- Who is treated as data before being treated as human?
- Who can stop the machine?

Closing line:

Before every system, there is a person. Before every dataset, there is a life. Before every optimization, there is a sacrifice.

UI Requirements

- Use a contemplative visual treatment
- Could be a centered litany-style section
- Strong spacing
- Must be readable as a standalone moral test

Acceptance Criteria

- Section is easy to quote and copy
- It gives the site a human, embodied emotional layer
- It balances the higher theological argument with practical moral imagination

Feature 8: Add Reader Paths

Location

Near top, after reading modes, or near the Doctrine Map.

Purpose

Help different users enter the material.

Required Paths

For Builders

Prompt:

How should I design AI without degrading human agency?

Target section:

Governance Gates

For Policymakers

Prompt:

What must be regulated?

Target section:

Five Red Lines / Common Good / Accountability

For Workers

Prompt:

What protections does dignity require?

Target section:

Truth, Work, Freedom

For Theologians

Prompt:

How does Catholic Social Doctrine extend into AI?

Target section:

Foundations & Principles

For Investors

Prompt:

Which AI companies build toward Jerusalem, not Babel?

Target section:

Babel vs. Jerusalem / Doctrine Map

UI Requirements

- Use cards or compact links
- Each card should jump to relevant section
- Do not over-design
- Maintain solemn editorial style

Acceptance Criteria

- User can self-identify and jump into the right material
 - Each path uses working anchors
 - Mobile layout remains clean
-

Feature 9: Upgrade Chapter Headers

Current Issue

Chapter sections are readable but should be more orienting.

Required Header Format

Each chapter header should include:

- Roman numeral
- Latin title
- English title
- Paragraph range
- One-sentence summary
- Key quote

Example

Chapter III:

Roman numeral:

III

Latin title:

De technologia et dominatione

English title:

Technology & Dominance

Paragraph range:

nn. 90–130

Summary:

AI becomes dangerous when technical power claims the right to govern human life.

Key quote:

Technology is never neutral.

Required Chapter Summaries

Introduction

The age of AI forces a choice between Babel — domination through efficiency — and Jerusalem — shared responsibility for human dignity.

Key quote:

Technology is never neutral; it takes on the characteristics of those who devise and finance it.

Chapter I

Catholic Social Doctrine is a living method of discernment, not a museum of fixed answers.

Key quote:

Social Doctrine is not a museum of answers, but a shared discernment between the eternal and the new.

Chapter II

The human person is made for communion and bears a dignity no machine may weigh.

Key quote:

Human dignity cannot be earned.

Chapter III

AI becomes dangerous when technical power claims the right to govern human life.

Key quote:

Humanity flourishes through its limitations, not in spite of them.

Chapter IV

Truth, work, and freedom must be protected from systems that exploit attention, labor, and vulnerability.

Key quote:

The pursuit of greater profit cannot justify the choices that systematically sacrifice human work.

Chapter V

A civilization of love requires accountability, restraint, justice, and human control over the tools of war.

Key quote:

Blame may never be collapsed into “the machine.”

Conclusion

What saves humanity is not an upgrade, but love descending into fragility.

Key quote:

What saves humanity is the divine love that descends into the most fragile point of our history and renews it from within.

Acceptance Criteria

- Each chapter gives the user orientation before the long text
 - Chapter summaries are visible but do not interrupt the ceremonial reading flow
 - Existing content remains intact
-

Feature 10: Add Quote Actions

Purpose

Turn strong quotes into reusable doctrine fragments.

Required Quotes

Add action buttons to major pull quotes:

- Copy quote
- Copy governance rule
- Jump to paragraph
- Copy anchor link

Required Governance Rule Translations

Quote:

Human dignity cannot be earned.

Governance rule:

Do not deploy AI systems that rank, exclude, punish, or exploit people based on predicted productivity, efficiency, vulnerability, or social value.

Quote:

Technology is never neutral.

Governance rule:

Require disclosure of the values, incentives, funding structures, and deployment context shaping the system.

Quote:

The pursuit of greater profit cannot justify the choices that systematically sacrifice human work.

Governance rule:

Require worker impact review, retraining plan, surveillance limits, and displacement mitigation before automation.

Quote:

Blame may never be collapsed into “the machine.”

Governance rule:

Require named accountable human owners, audit trails, appeal paths, and incident response protocols.

Technical Notes

- Use JavaScript clipboard API
- Provide visual feedback after copy
- Example feedback: “Copied”
- Buttons must be keyboard accessible
- If clipboard API fails, fallback gracefully

Acceptance Criteria

- Quote actions work on desktop and mobile
- Copy actions do not break layout
- No console errors
- Buttons are accessible by keyboard

Feature 11: Add Governance Export Area

Location

Near end of page, before conclusion or after Governance Gates.

Title

Export the Doctrine

Subtitle

Use Magnifica Humanitas as an AI governance skill, audit checklist, or agent instruction set.

Required Export Cards

SKILL.md

Description:

A governance / ethics agent skill using the five principles and verification gates.

Button:

Copy SKILL.md

principles.yaml

Description:

Structured principles: dignity, common good, subsidiarity, truth, and accountable human control.

Button:

Copy principles.yaml

audit_gates.yaml

Description:

Verification gates for reviewing high-stakes AI systems.

Button:

Copy audit_gates.yaml

agent_prompt.md

Description:

A prompt for an AI governance reviewer agent.

Button:

Copy agent_prompt.md

MVP Behavior

For now, copying predefined text to clipboard is enough. No file generation is required unless easy.

Acceptance Criteria

- User can copy at least SKILL.md content
 - Export area does not visually cheapen the site
 - Copy feedback works
 - Export content is well formatted
-

Feature 12: Add Sticky Doctrine Rail

Purpose

Improve orientation during long reading.

Desktop Behavior

Sticky side rail shows:

- Current chapter
- Paragraph range
- Active principle or theme
- Key question

Example:

Current chapter:

Technology & Dominance

Paragraph range:

nn. 90–130

Active theme:

Technocratic Paradigm

Key question:

Who decides what shall matter?

Mobile Behavior

Do not use sticky side rail on small screens. Replace with compact section progress / current chapter label.

Acceptance Criteria

- Sticky rail does not overlap content
 - Hidden or simplified on mobile
 - Current section updates if technically feasible
 - If current-section tracking is too heavy, static anchor rail is acceptable for MVP
-

Feature 13: Accessibility and Responsiveness Pass

Requirements

- All text must meet reasonable contrast standards
- Buttons must be keyboard navigable
- Anchor links must work
- Focus states must be visible
- Copy buttons must have accessible labels
- Avoid tiny text on mobile
- Ensure line length does not become too wide on desktop
- Ensure no horizontal scrolling on mobile
- Respect reduced-motion settings if smooth scrolling or animations are added

Acceptance Criteria

- Site is usable on desktop, tablet, and mobile
 - No critical Lighthouse accessibility failures
 - Navigation is keyboard-accessible
 - Copy buttons announce state visually at minimum
-

Innovative Approach:

14. Product Vision Upgrade

Do not simply “add sections.”

The deeper goal is to make the site feel like a **moral operating system** for the AI age.

The user should leave with the sense that this is not only an encyclical, but a tool for judgment.

Target feeling:

A cathedral console for judging machines.

Design keywords:

- Ecclesial
- Archival
- Constitutional
- Judicial
- Liturgical
- Severe
- Human
- Anti-SaaS
- Anti-dashboard cliché

Avoid:

- Startup gradients
 - Neon AI aesthetics
 - Cute cards
 - Generic “ethics checklist” vibes
 - Overly colorful tags
 - Excessive animations
 - Dashboard chrome
-

15. Suggested Information Architecture

Recommended order:

1. Cover / Hero
2. Executive Summary
3. Reading Modes
4. Reader Paths
5. Table of Contents
6. Introduction
7. Chapter I
8. Chapter II
9. Five Principles for the Digital Age
10. Doctrine Map
11. Chapter III
12. Chapter IV
13. Governance Gates
14. Five Red Lines
15. The Person Test

- 16. Chapter V
- 17. Export the Doctrine
- 18. Conclusion

Alternative:

If the agent wants to keep the encyclical text uninterrupted, place all new operational sections after the Table of Contents and before the full text as a “Doctrine Console.”

16. Implementation Notes for Coding Agent

Assumption

The site appears to be a static editorial page hosted on GitHub Pages. If the implementation is plain HTML/CSS/JS, make changes directly in:

- `index.html`
- main stylesheet
- main JavaScript file if present

If the site uses a framework, preserve the existing component structure and create reusable components.

Suggested Components

- `ExecutiveSummary`
- `ReadingModes`
- `ReaderPaths`
- `PrincipleCards`
- `DoctrineMap`
- `GovernanceGates`
- `RedLines`
- `PersonTest`
- `QuoteActions`
- `ExportDoctrine`
- `StickyDoctrineRail`
- `ChapterHeader`

Anchor IDs

Use stable IDs:

- `#executive-summary`
- `#reading-modes`
- `#reader-paths`

- #principles
- #doctrine-map
- #governance-gates
- #red-lines
- #person-test
- #export-doctrine

JavaScript Requirements

Implement:

- Smooth scrolling
- Clipboard copy for quotes
- Clipboard copy for governance exports
- Copy success state
- Optional active-section detection for sticky rail

CSS Requirements

Preserve current visual identity.

Add:

- Card styles
- Gate styles
- Litany / quotation styles
- Responsive grid utilities
- Accessible button states
- Mobile stacking behavior
- Print-friendly layout if easy

17. QA Checklist

Before final delivery, verify:

- All new anchor links work
- Existing table of contents still works
- Dropdown navigation still works
- No existing encyclical text is removed accidentally
- Mobile layout is readable
- No horizontal scrolling on mobile
- Copy buttons work
- Copy fallback does not crash
- Sticky rail does not cover content

- Page loads without console errors
 - Accessibility focus states are visible
 - All new sections preserve the site's tone
 - The site still feels like Magnifica Humanitas, not a generic AI ethics template
-

18. Definition of Done

The task is complete when:

- The site includes the new summary, doctrine, and governance layers
 - Users can choose a reading path
 - Users can understand the five principles quickly
 - Users can apply the governance gates to an AI system
 - Users can copy quotes or governance rules
 - Users can export at least SKILL.md content
 - Mobile experience is clean
 - Existing encyclical content is preserved
 - The final page feels more operational without losing its sacred / ceremonial atmosphere
-

What I should remember

- Preserve the encyclical's solemn authority.
- Add usability without making it feel like SaaS.
- Make the doctrine scannable, applicable, and exportable.
- The strongest UX idea is the move from document → doctrine map → governance engine.
- The final site should feel like a moral architecture for judging AI systems.